

BORDER-AI

Edge-Intelligent, Multi-Camera Video Analytics & Event Intelligence Platform for Existing Border CCTV Infrastructure

SIH26187 — Ministry of Home Affairs | Sashastra Seema Bal (SSB), Police II Division | Category: Software | Theme: Smart Automation

1. Introduction & Summary

BORDER-AI is a software platform that converts existing, conventional CCTV infrastructure at Border Outposts, check posts and border roads into an intelligent, autonomous video analytics and event-intelligence system. Rather than replacing deployed cameras with expensive smart-camera hardware, the platform adds intelligence entirely through software: it ingests existing video streams, detects and tracks people and vehicles, understands behaviour over time, correlates activity across multiple cameras, and produces explainable, evidence-backed alerts for a human operator.

The core idea is to move beyond a single-frame object detector and build a complete event-intelligence pipeline: Detect → Track → Understand → Correlate → Score → Alert → Explain.

2. Background / Problem

Border surveillance today relies on conventional CCTV cameras that provide continuous recording and live video, but the intelligence layer is still human: personnel must continuously watch multiple feeds to catch intrusions, suspicious movement or unattended objects. The problem statement calls out specific desired capabilities — facial recognition, automatic number-plate recognition, intrusion detection, object tracking and activity recognition — but deploying dedicated hardware for each of these is costly and difficult, particularly in remote border areas with unreliable connectivity.

The opportunity is therefore not to build “CCTV → YOLO → alert”, which is the obvious and commonly attempted solution, but a software platform that layers detection, tracking, temporal reasoning, multi-camera correlation and explainable risk-scoring on top of the video an organisation already has.

3. Motivation

- Border areas are remote, with limited bandwidth and unreliable connectivity, ruling out cloud-only architectures.
- Replacing existing CCTV with dedicated smart cameras for every required function (FRS, ANPR, intrusion, tracking, activity recognition) is prohibitively expensive at national scale.
- Operators cannot reliably monitor dozens of simultaneous feeds; fatigue and attention limits are a real operational risk.
- A single-frame alert (“person detected”) is not actionable; operators need context — how long, where, moving toward what, correlated with which other camera.
- Government and defence deployments require explainability, auditability and privacy safeguards, not just raw detection accuracy.

4. Objectives

- Ingest existing CCTV/RTSP streams without requiring new hardware at the camera end.
- Detect and track people and vehicles in real time using an edge-deployable detector and tracker.
- Detect restricted-zone intrusion and loitering with temporal confirmation to minimise false alarms.
- Detect abandoned objects and unusual multi-person coordination.
- Read vehicle number plates (ANPR) with a confidence score rather than a binary claim of correctness.
- Support an optional, privacy-conscious authorized-watchlist face-matching module, not a core identity system.
- Correlate events across multiple cameras into a single cross-camera event timeline.
- Produce an explainable, component-wise risk score for every event, not just a single opaque number.

- Operate under low-bandwidth or intermittently offline conditions, synchronising only event metadata and short evidence clips when connectivity is available.

5. Proposed Solution

Core idea: BORDER-AI is an edge-AI surveillance platform that treats every existing camera as an intelligent sensor rather than a passive recorder. Detection is only the first stage; the differentiating value comes from tracking objects over time, reasoning about their behaviour, correlating events across cameras, and presenting operators with an explainable, evidence-backed alert instead of a raw detection.

A restricted-zone crossing is not reported as an isolated frame-level detection. It is combined with tracking persistence, dwell time, movement direction and, where relevant, activity at other cameras, before it is escalated to the operator as a single high-risk event with a stated reason.

6. Proposed System Architecture

End-to-end pipeline

1. Existing CCTV cameras stream video (RTSP) into the platform — no camera replacement required.
2. A video ingestion layer (OpenCV / GStreamer) manages streams and performs frame sampling.
3. A fast object-detection model (YOLO-family) runs continuously, detecting people and vehicles.
4. Detections are handed to a multi-object tracker (ByteTrack / BoT-SORT), which assigns persistent track IDs.
5. Person tracks feed a behaviour-analysis stage (zone intrusion, loitering); vehicle tracks feed vehicle analytics and ANPR.
6. A temporal AI engine converts sequences of low-level detections into higher-level activities and flags anomalies.
7. A multi-camera fusion layer links related tracks and events across different cameras into one timeline.
8. A risk engine scores each event using explainable, additive rules and assigns a severity band.
9. An alert-generation layer pushes evidence-backed alerts to the command dashboard and notification channels.

A second, lighter-weight tier is triggered only when an event occurs: a fast detector runs continuously (Tier 1) while a more detailed analysis model (Tier 2) is invoked only on candidate events, reducing continuous compute load — an important engineering argument for constrained edge hardware.

7. AI / ML Modules

7.1 Object detection

Detects person, vehicle, motorcycle, car, truck, bus and other relevant object classes using a modern YOLO-family or comparably efficient detector.

7.2 Object tracking

ByteTrack or BoT-SORT assigns each detected object a persistent track ID, turning isolated detections (“there is a person”) into a continuous record (“this is the same person who entered 40 seconds ago and is now moving east”), which is what enables temporal reasoning.

7.3 Intrusion detection

Operators draw virtual restricted zones on each camera view. A crossing is not alerted from a single frame; it requires zone crossing plus tracking persistence plus temporal confirmation, which materially reduces false alarms.

7.4 Loitering detection

Dwell time inside a zone is classified on a configurable scale (for example: under 5 seconds is normal, 30 seconds is monitored, 90 seconds or a configurable threshold triggers an alert), combined with zone sensitivity, movement pattern and time-of-day context to compute a contribution to the risk score.

7.5 Abandoned object detection

When a tracked person places an object and then leaves it behind, the system verifies persistence over time before raising an alert (e.g. “possible abandoned object, Camera BOP-07, 72 seconds, restricted area”), again avoiding single-frame false positives.

7.6 Vehicle analytics

Tracks vehicle entry, exit, direction, approximate speed (where calibration permits), loitering and restricted-zone entry — useful at check posts, border roads and entry points.

7.7 Automatic number-plate recognition (ANPR)

Pipeline: vehicle detection → vehicle tracking → plate detection → plate cropping → image enhancement → OCR → plate text → validation. Output is reported as an OCR confidence score (e.g. 91%) with operator verification, rather than a claim of 100% accuracy.

7.8 Face recognition — authorized watchlist only

Rather than making facial recognition a core identity system (which raises privacy, false-match, dataset and legal concerns), it is scoped as an optional watchlist-comparison module: face detected → quality check → embedding → authorized-watchlist comparison → similarity score → mandatory human verification. Results are reported as a “possible match”, never a definitive identification.

7.9 Temporal activity intelligence

The system's central differentiator: instead of reasoning frame-by-frame, it represents an event as a sequence — object, position, velocity, direction, dwell duration, zone and prior related events — and combines related low-level detections into a single higher-level narrative, e.g. “person entered the restricted zone and remained for approximately 94 seconds before moving toward the boundary,” instead of five disconnected alerts.

7.10 Multi-camera correlation

Related tracks and events across different cameras are linked into a cross-camera event timeline, giving the operator the complete journey of a person or vehicle across the deployment rather than isolated, camera-by-camera detections.

7.11 Explainable risk-scoring engine

Every event receives an additive, component-wise score rather than an opaque number, so operators can see exactly why an alert was raised.

Event / factor	Score contribution
Person detected	+5
Restricted-zone entry	+30
Loitering	+20
Movement toward restricted boundary	+20
Multiple-person coordination	+15
Abandoned object	+30

Resulting severity bands: 0–30 Low, 31–60 Medium, 61–80 High, 81–100 Critical. Every alert is displayed with its full breakdown (for example: Risk = 87 — +30 restricted zone, +20 loitering, +20 boundary direction, +17 multi-event correlation) so the reasoning is visible, not just the conclusion.

8. Tech Stack

Layer	Technology	Purpose
Computer vision	OpenCV, PyTorch, Ultralytics YOLO-family, ONNX Runtime	Detection, video processing, optimized inference

Tracking	ByteTrack / BoT-SORT	Persistent multi-object tracking
OCR	OCR framework (e.g. PaddleOCR/Tesseract-class)	ANPR plate-text extraction
Backend	FastAPI	REST APIs, event/risk services
Frontend	React.js	Command dashboard, camera grid, map view
Database	PostgreSQL (+ vector extension optional)	Cameras, events, alerts, tracks, users, audit logs
Storage	Object storage	Evidence snapshots and short event clips only
Deployment	Docker, microservices	stream / ai / tracking / event / alert / api services
Edge targets	NVIDIA Jetson-class device, GPU workstation, CPU fallback	On-site inference without cloud dependency

9. Data Required / Dataset Strategy

A large, real SSB border-CCTV dataset will not realistically be available, so a layered dataset strategy is required:

- Layer 1 — Public datasets for general person/vehicle detection, tracking and activity recognition.
- Layer 2 — Synthetic, self-recorded scenarios: normal walking, restricted-zone crossing, loitering, abandoned object, vehicle movement, group movement.
- Layer 3 — Data augmentation for realistic CCTV conditions: night, low brightness, blur, rain, compression artefacts, low resolution, occlusion.
- Layer 4 — Domain adaptation so final demo data resembles fixed-position outdoor surveillance cameras rather than curated benchmark footage.

For low-light conditions, an enhancement stage (low-light frame → enhancement → detection → tracking) can be evaluated, but any claim of improved detection accuracy should be benchmarked rather than assumed.

10. Key Innovations & Differentiators

- Existing-CCTV-first: no requirement to replace deployed cameras with specialised smart hardware.
- Edge-first AI: inference runs locally, so the system functions under limited or intermittent connectivity.
- Temporal intelligence: events are understood as sequences over time, not isolated frames.
- Multi-camera correlation: a person or vehicle's journey is tracked across the entire camera network, not per-camera.
- Explainable risk scoring: every alert shows the factors that produced its score.
- Adaptive, operator-defined zones: restricted areas are configurable per camera, not hard-coded.
- False-alarm reduction via temporal confirmation before any critical alert is raised.
- Human-in-the-loop design: AI assists the operator's decision, it does not replace it.
- Evidence-based alerts: every alert carries a snapshot, timestamp, camera, track ID, reason, confidence and event timeline.
- Offline / low-bandwidth operation: only event metadata and short evidence clips synchronise when connectivity allows.
- Event replay: operators can view 10 seconds before and 20 seconds after any alert, with bounding boxes and track ID overlaid.
- Natural-language event search: an LLM translates operator questions (“show all intrusion events from Camera 3 today”) into controlled queries over structured event data — it never controls the surveillance system directly.

- AI-generated incident summaries: a short, human-readable narrative is generated for each event from its structured data.

11. Edge AI & Low-Bandwidth Operation

Because the problem statement explicitly concerns remote border areas and deployment cost, the architecture avoids a cloud-GPU-dependent design (which would introduce bandwidth dependency, latency, privacy exposure and recurring cost). Instead, inference runs at the edge: camera → local edge AI → event metadata → low-bandwidth server → command dashboard. Full video need not be transmitted continuously — only relevant clips and metadata for confirmed events, which is itself a strong, measurable engineering argument (bandwidth saved versus continuous video transmission).

12. Security, Privacy & Explainability

Security

- Role-based access control: Admin, Operator, Supervisor, Auditor.
- TLS for communications and encryption of storage where appropriate.
- Full audit logging of who did what, when, and from where.

Privacy-by-design

- Configurable masking of faces and plates unless an authorized workflow requires access.
- Defined data-retention policies — footage is not retained indefinitely by default.
- Mandatory human verification before any face-match or plate-read is treated as confirmed.

Explainability

No alert is presented as a bare conclusion (“AI detected threat”). Every alert is shown with the specific evidence and score components that produced it, which is both more trustworthy for an operator and more defensible to judges and to real deployment stakeholders.

13. Impact / Benefits

- Reduces continuous manual monitoring load across dozens of simultaneous camera feeds.
- Converts existing CCTV investment into an intelligent system without hardware replacement cost.
- Faster, more consistent detection of intrusion, loitering and abandoned objects than human attention alone.
- Cross-camera event timelines give a complete picture of an incident, not fragments.
- Explainable, evidence-backed alerts improve operator trust and reduce time to decision.
- Functions under poor or intermittent connectivity, which is realistic for remote border terrain.
- Demonstrable bandwidth savings versus continuous video transmission.

14. Key USP

Core innovation: the platform's value is not any single detector — it is the complete event-intelligence pipeline: detection, persistent tracking, temporal behaviour analysis, multi-camera correlation, explainable risk scoring and evidence-based alerting, deployed at the edge so it keeps working when connectivity does not. This directly answers the most likely judge question (“why is this better than YOLO?”): YOLO is one component among many, not the product.

15. Demo Scenario & Presentation Strategy

Rather than demonstrating many partially-built features, the recommended approach is one complete, convincing scenario, shown end-to-end: a normal CCTV feed → a person is detected → the person enters a restricted zone → the system flags intrusion → the person loiters and the risk score rises → the person moves toward the boundary → a second camera detects the same track and the events are correlated → the dashboard raises a single high-risk event → the operator opens it to see video, timeline, snapshot, tracking and the explainable score → the operator asks the natural-language assistant to show all high-risk events from the last hour.

Recommended build priority: get intrusion + tracking + loitering + alert + dashboard working completely first. Only after that pipeline is solid end-to-end should ANPR, multi-camera correlation, edge/offline mode and the natural-language assistant be added as differentiators.

16. Expected Outcome

A deployable proof-of-concept platform that ingests existing CCTV streams and outputs explainable, evidence-backed, cross-camera security events to a command dashboard, operating reliably under the low-bandwidth, remote conditions typical of border deployments — without requiring camera hardware replacement.

Existing CCTV + Edge AI + Temporal Intelligence + Multi-Camera Correlation + Explainable Risk + Low-Bandwidth Operation